

Pure_TensorFlow

September 29, 2020

```
[1]: import tensorflow as tf
import datetime
import numpy as np
import matplotlib.pyplot as plt

tf.__version__
tf.random.set_seed(44)
np.random.seed(44)
```

```
[2]: class Dense(tf.Module):
    def __init__(self, input_size, output_size, name=None,
        ↪activation_function=tf.nn.relu):
        super().__init__(name=name)
        self.activation_function = activation_function
        self.weights = tf.Variable(tf.random.normal([input_size, output_size]),
        ↪name='weights')
        self.bias = tf.Variable(tf.random.normal([output_size]), name='bias')

    @tf.function
    def __call__(self, x):
        y = tf.matmul(x, self.weights) + self.bias
        return self.activation_function(y)

class NeuralNetwork(tf.Module):
    def __init__(self, input_size, layers, name=None):
        super(NeuralNetwork, self).__init__(name=name)
        self.layers = []
        l = 0
        with self.name_scope:
            for size in layers:
                if l == len(layers) - 1:
                    self.layers.append(Dense(input_size=input_size,
        ↪output_size=size, name=f'dense_{l}', activation_function = tf.nn.sigmoid))
                    break

                self.layers.append(Dense(input_size=input_size,
        ↪output_size=size, name=f'dense_{l}'))
```

```

        input_size = size
        l += 1

    @tf.Module.with_name_scope
    def __call__(self, x):
        for layer in self.layers:
            x = layer(x)
        return x

```

```

[3]: def loss(target_y, predicted_y):
      return tf.reduce_mean(tf.square(target_y - predicted_y))

```

```

[4]: def accuracy(target_y, predicted_y):
      return (100.0 / len(target_y)) * sum(target_y == tf.math.round(predicted_y).
      ↪numpy())

```

```

[5]: inputs = np.array([[0.0,0.0],[0.0,1.0],[1.0,0.0],[1.0,1.0]], dtype=np.float32)
      expected_output = np.array([[0.0],[1.0],[1.0],[0.0]], dtype=np.float32)

```

```

[6]: def train(model, x, y, learning_rate):
      with tf.GradientTape(persistent=True) as t:
          predicted_y = model(x)
          current_loss = loss(y, predicted_y)
          for layer in reversed(model.layers):
              dw, db = t.gradient(current_loss, [layer.weights, layer.bias])
              layer.weights.assign_sub(learning_rate * dw)
              layer.bias.assign_sub(learning_rate * db)

```

```

[7]: model = NeuralNetwork(input_size = 2, layers = [6, 6, 6, 1],
      ↪name="MyNeuralNetwork")

```

```

epochs = range(100)
losses = []
accuracies = []

def training_loop(model, x, y):
    for epoch in epochs:
        train(model, x, y, learning_rate=0.1)
        predicted_y = model(x)
        current_loss = loss(y, predicted_y)
        current_accuracy = accuracy(y, predicted_y)
        accuracies.append(current_accuracy)
        losses.append(current_loss)
        if epoch % 100 == 0:
            print(f"Epoch {epoch}: loss={current_loss},
            ↪accuracy={current_accuracy[0]}%")

```

```
[8]: training_loop(model, inputs, expected_output)
```

WARNING:tensorflow:Calling GradientTape.gradient on a persistent tape inside its context is significantly less efficient than calling it outside the context (it causes the gradient ops to be recorded on the tape, leading to increased CPU and memory usage). Only call GradientTape.gradient inside the context if you actually want to trace the gradient in order to compute higher order derivatives.

Epoch 0: loss=0.2728033661842346, accuracy=50.0%

```
[9]: fig, axes = plt.subplots(1, 2, figsize=(20, 10))
    axes[0].plot(losses)
    axes[0].set_title('Loss')
    axes[1].plot(accuracies)
    axes[1].set_title('Accuracy')
    fig.suptitle('Predicting XOR')
```

```
[9]: Text(0.5, 0.98, 'Predicting XOR')
```

