

# CNN - Exercise

In the tutorial Standing on the Shoulders of Giants you investigated a computer vision problem of your choice by applying transfer learning to a series of input images of your choice. For this exercise, you should re-use the weights for the convolutional layers and train the other layers of the model using your own input images and evaluate the results. Hand in your Jupyter notebook which also functions as a research report. You describe the context of the problem at hand, the methods, the results and your conclusion.

## Preparation

In this section useful libraries are imported which are used in most data science projects.

```
import os
import sys

# sets the path to the home directory of this repository so other modules can be imported.
project_path = os.getcwd()
root_path = os.path.split(os.path.split(os.getcwd())[0])[0]
assert root_path.endswith("Fontys-ADS"), "The root path does not end with Fontys-ADS: " +
root_path
sys.path.insert(0, root_path)

import numpy as np
import tensorflow as tf
print(tf.__version__)
# set the seed for reproducible results.
np.random.seed(56)
tf.random.set_seed(56)

# optionally, set TensorFlow to use the GPU with all available memory.
physical_devices = tf.config.experimental.list_physical_devices('GPU')
tf.config.experimental.set_memory_growth(physical_devices[0], True)
```

2.3.0

## Data collection

For this assignment, I have used the cats\_vs\_dogs image dataset neatly provided by the TensorFlow datasets API.

```
import pandas as pd
import tensorflow_datasets as tfds
ds, info = tfds.load('cats_vs_dogs', split='train', with_info=True)
assert isinstance(ds, tf.data.Dataset)
```

# Preparing the data

explain how the data is prepared

```
from datasets.base_image_dataset import ImageDatasetBase

# the dataset class
class CatsVsDogsDataset(ImageDatasetBase):
    def __init__(self, ds, batch_size, img_height, img_width, data_size, train_percentage,
validation_percentage, test_percentage):
        super().__init__(batch_size, img_height, img_width)

        # sets the data.
        self.data = ds.map(self.process_frame, num_parallel_calls=tf.data.experimental.AUT
OTUNE)

        # shuffles the dataset
        self.shuffle(256)

        # splits the data into train, validation, and test datasets.
        self.split_data_to_train_val_test(self.data, train_percentage, validation_percenta
ge, test_percentage, data_size / batch_size)

    def process_frame(self, ds):
        # convert image to 0..1
        img = tf.image.convert_image_dtype(ds['image'], tf.float32)

        # resize the image to the desired size.
        return tf.image.resize(img, [self.img_height, self.img_width]), ds['label']
```

```
data_size = 23262
img_height = 224
img_width = 224
img_channels = 3
classes = 2

batch_size = 64
train_percentage = 0.6
validation_percentage = 0.2
test_percentage = 0.2
dataset = CatsVsDogsDataset(ds, batch_size, img_height, img_width, data_size, train_percen
tage, validation_percentage, test_percentage)
```


train: 218 validation: 72 test: 72

## Exploratory Data Analysis

Explore the data to gain insights on possible features

```
import matplotlib.pyplot as plt
class_names = ['Cat', 'Dog']
for img, label in dataset.train_ds.take(1):
    print(len(img.numpy()))
    fig, ax = plt.subplots(1, 1)
    ax.imshow(img[0])
    ax.set_title(class_names[label[0].numpy()])
```

64



## Modelling

Apply ML/DL models

```

from models.base_model import ModelBase
from tensorflow.keras import Model
from tensorflow.keras.layers import Input, Conv2D, MaxPooling2D, Dropout, Flatten, Dense
from tensorflow.keras.activations import elu

class SuperCatDogClassifier(ModelBase):
    def __init__(self, img_height, img_width, img_channels, classes, gpu_initialized=False
, training=False, limit=5000):
        super().__init__(gpu_initialized, training, limit)
        # the name for the model.
        self.name = 'SuperCatDogClassifier'

        # set img dimensions.
        self.img_height = img_height
        self.img_width = img_width
        self.img_channels = img_channels

        # set the classes.
        self.classes = classes

    def predict(self, X):
        prediction = self.model.predict(X, verbose=1)
        return prediction

    def fit(self, training, callbacks, epochs, validation, validation_steps, steps_per_epoch):
        return self.model.fit(
            training,
            callbacks=callbacks,
            epochs=epochs,
            validation_data=validation,
            validation_steps=validation_steps,
            steps_per_epoch=steps_per_epoch, verbose=1)

    def compile(self, optimizer='adam', loss='mse', metrics=['mse'], loss_weights=[1.0], show_summary=False):
        self.inputs = Input((self.img_height, self.img_width, self.img_channels))

        c1 = Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(self.inputs)
        p1 = MaxPooling2D((2, 2))(c1)

        c9 = Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(p1)
        p2 = MaxPooling2D((2, 2))(c9)

        c2 = Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(p2)
        p3 = MaxPooling2D((2, 2))(c2)

        f = Flatten()(p3)
        d1 = Dense(128, activation='relu', kernel_initializer='he_uniform')(f)

        self.outputs = Dense(1, activation='sigmoid')(d1)

        self.model = Model(inputs=self.inputs, outputs=self.outputs, name=self.name)

```

```
        self.model.compile(optimizer=optimizer, loss=loss, metrics=metrics, loss_weights=loss_weights)

        if show_summary:
            self.model.summary()
```

```
model = SuperCatDogClassifier(img_height, img_width, img_channels, classes, training=True,
                             gpu_initialized=True)
```

```

from tensorflow.keras.optimizers import Adam
from tensorflow.keras.callbacks import EarlyStopping, TensorBoard
import datetime

epochs = 100
model.compile(optimizer=Adam(lr = 1e-4), loss='binary_crossentropy', metrics=['accuracy'],
show_summary=True)

# current time
current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")

# create logging
log_dir = os.path.join(project_path, f'logs\\{model.name}\\{current_time}')

# create all callbacks
callbacks = [
    EarlyStopping(patience=10, monitor='val_loss'),
    TensorBoard(log_dir=log_dir, profile_batch=0)
]

```

Model: "SuperCatDogClassifier";

Layer (type)	Output Shape	Param #
=====		
input_1 (InputLayer)	[(None, 224, 224, 3)]	0
-----		
conv2d (Conv2D)	(None, 224, 224, 32)	896
-----		
max_pooling2d (MaxPooling2D)	(None, 112, 112, 32)	0
-----		
conv2d_1 (Conv2D)	(None, 112, 112, 32)	9248
-----		
max_pooling2d_1 (MaxPooling2)	(None, 56, 56, 32)	0
-----		
conv2d_2 (Conv2D)	(None, 56, 56, 32)	9248
-----		
max_pooling2d_2 (MaxPooling2)	(None, 28, 28, 32)	0
-----		
flatten (Flatten)	(None, 25088)	0
-----		
dense (Dense)	(None, 128)	3211392
-----		
dense_1 (Dense)	(None, 1)	129
=====		
Total params: 3,230,913		
Trainable params: 3,230,913		
Non-trainable params: 0		
-----		

```
%load_ext tensorboard
```

```
# fit the model using the training data
```

```
results1 = model.fit(  
    training=dataset.train_ds,  
    callbacks=callbacks,  
    epochs=epochs,  
    validation=dataset.val_ds,  
    validation_steps=dataset.val_size,  
    steps_per_epoch=dataset.train_size)
```

```
# save the weights of the model
```

```
weights_path = os.path.join(project_path, f'models\{model.name}_trained_model_weights')  
model.save_weights(weights_path)
```

Epoch 1/100

218/218 [=====] - 33s 153ms/step - loss: 0.6393 - accuracy: 0.6262 - val\_loss: 0.5740 - val\_accuracy: 0.7064

Epoch 2/100

218/218 [=====] - 12s 54ms/step - loss: 0.5501 - accuracy: 0.7183 - val\_loss: 0.5222 - val\_accuracy: 0.7422

Epoch 3/100

218/218 [=====] - 11s 49ms/step - loss: 0.5018 - accuracy: 0.7577 - val\_loss: 0.4931 - val\_accuracy: 0.7598

Epoch 4/100

218/218 [=====] - 11s 50ms/step - loss: 0.4662 - accuracy: 0.7795 - val\_loss: 0.4763 - val\_accuracy: 0.7652

Epoch 5/100

218/218 [=====] - 11s 49ms/step - loss: 0.4353 - accuracy: 0.7992 - val\_loss: 0.4664 - val\_accuracy: 0.7734

Epoch 6/100

218/218 [=====] - 11s 49ms/step - loss: 0.4103 - accuracy: 0.8154 - val\_loss: 0.4597 - val\_accuracy: 0.7812

Epoch 7/100

218/218 [=====] - 11s 48ms/step - loss: 0.3876 - accuracy: 0.8293 - val\_loss: 0.4547 - val\_accuracy: 0.7878

Epoch 8/100

218/218 [=====] - 11s 48ms/step - loss: 0.3663 - accuracy: 0.8413 - val\_loss: 0.4522 - val\_accuracy: 0.7893

Epoch 9/100

218/218 [=====] - 11s 50ms/step - loss: 0.3455 - accuracy: 0.8534 - val\_loss: 0.4513 - val\_accuracy: 0.7884

Epoch 10/100

218/218 [=====] - 11s 51ms/step - loss: 0.3257 - accuracy: 0.8655 - val\_loss: 0.4531 - val\_accuracy: 0.7880

Epoch 11/100

218/218 [=====] - 10s 47ms/step - loss: 0.3069 - accuracy: 0.8743 - val\_loss: 0.4577 - val\_accuracy: 0.7888

Epoch 12/100

218/218 [=====] - 10s 47ms/step - loss: 0.2871 - accuracy: 0.8845 - val\_loss: 0.4587 - val\_accuracy: 0.7938

Epoch 13/100

218/218 [=====] - 10s 47ms/step - loss: 0.2682 - accuracy: 0.8954 - val\_loss: 0.4626 - val\_accuracy: 0.7973

Epoch 14/100

218/218 [=====] - 10s 47ms/step - loss: 0.2509 - accuracy: 0.9036 - val\_loss: 0.4692 - val\_accuracy: 0.7949

Epoch 15/100

218/218 [=====] - 10s 47ms/step - loss: 0.2323 - accuracy: 0.9135 - val\_loss: 0.4781 - val\_accuracy: 0.7925

Epoch 16/100

218/218 [=====] - 11s 48ms/step - loss: 0.2187 - accuracy: 0.9200 - val\_loss: 0.5031 - val\_accuracy: 0.7849

Epoch 17/100

218/218 [=====] - 11s 49ms/step - loss: 0.2123 - accuracy: 0.9204 - val\_loss: 0.5655 - val\_accuracy: 0.7635

Epoch 18/100

218/218 [=====] - 11s 50ms/step - loss: 0.2190 - accuracy: 0.9127 - val\_loss: 0.5334 - val\_accuracy: 0.7763

Epoch 19/100

218/218 [=====] - 11s 49ms/step - loss: 0.2285 - accuracy: 0.9042 - val\_loss: 0.6366 - val\_accuracy: 0.7520



# Evaluation

Evaluation of the model performance

```
# re initialize the model.
model.training = False
model.compile(optimizer=Adam(lr = 1e-4), loss='binary_crossentropy', metrics=['accuracy'],
show_summary=False)
model.load_weights(weights_path)

print('\n# Evaluate on test data')
result = model.evaluate(dataset.actual_test_ds)
print('test loss, test acc:', result)
res = dict(zip(model.get_metric_names(), result))
print(res)
```

# Evaluate on test data

```
62/62 [=====] - 4s 61ms/step - loss: 0.6775 - accuracy: 0.7384
test loss, test acc: [0.6774950623512268, 0.7384072542190552]
{&#39;loss&#39;: 0.6774950623512268, &#39;accuracy&#39;: 0.7384072542190552}
```

```
image_batch, label_batch = next(iter(dataset.actual_test_ds))
y_pred = model.predict((image_batch, label_batch))
```

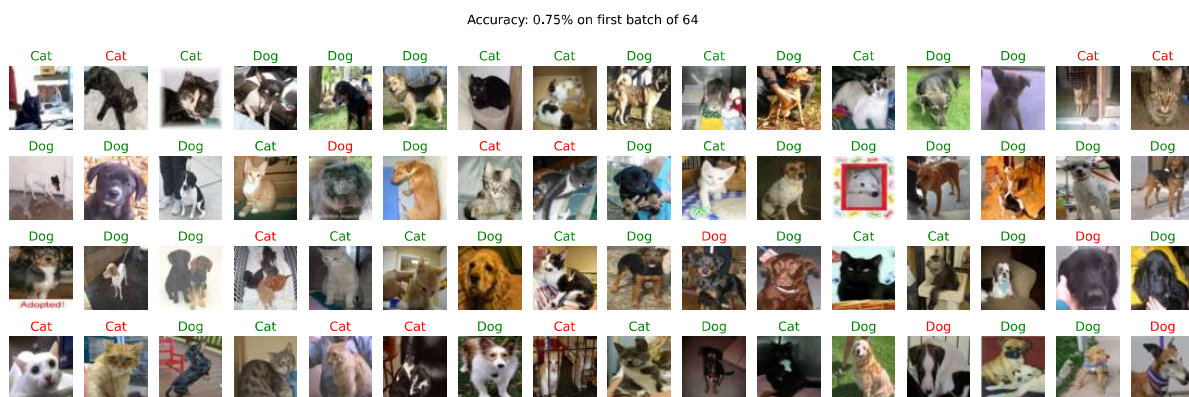
```
2/2 [=====] - 0s 5ms/step
```

```

predictions = [round(yhat[0]) for yhat in y_pred]
test_accuracy = sum(predictions == label_batch.numpy()) / len(predictions)
fig = plt.figure(figsize=(20, 6))
for i in range(batch_size):
    ax = plt.subplot(4, 16, i + 1)
    plt.imshow(image_batch[i])
    label = label_batch[i]
    pred = predictions[i]
    color = 'r'
    if label == pred:
        color = 'g'
    fontdict = { 'color': color }
    plt.title(class_names[label], fontdict = fontdict)
    plt.axis("off")
fig.suptitle(f'Accuracy: {test_accuracy:.2f}% on first batch of {batch_size}')

```

Text(0.5, 0.98, &#39;Accuracy: 0.75% on first batch of 64&#39;)



## Results

The results of my cat vs dog classification model was not the best, but 73% on a fairly simple architecture is quite good.

## Standing on the shoulder of giants

In this part of the notebook I will build a model using a pre trained base model as the feature extractor.

```

class SuperCatDogClassifierV2(ModelBase):
    def __init__(self, img_height, img_width, img_channels, classes, gpu_initialized=False
, training=False, limit=5000):
        super().__init__(gpu_initialized, training, limit)
        # the name for the model.
        self.name = 'SuperCatDogClassifierV2'

        # set img dimensions.
        self.img_height = img_height
        self.img_width = img_width
        self.img_channels = img_channels

        # set the classes.
        self.classes = classes

    def predict(self, X):
        prediction = self.model.predict(X, verbose=1)
        return prediction

    def fit(self, training, callbacks, epochs, validation, validation_steps, steps_per_epoch):
        return self.model.fit(
            training,
            callbacks=callbacks,
            epochs=epochs,
            validation_data=validation,
            validation_steps=validation_steps,
            steps_per_epoch=steps_per_epoch, verbose=1)

    def compile(self, optimizer='adam', loss='mse', metrics=['mse'], loss_weights=[1.0], show_summary=False):
        self.inputs = Input((self.img_height, self.img_width, self.img_channels))

        x = tf.keras.layers.experimental.preprocessing.Rescaling(2, offset=-1)(self.inputs)

        base_model = tf.keras.applications.MobileNetV2(input_shape=(self.img_height, self.img_width, self.img_channels), include_top=False, weights='imagenet')
        base_model.trainable = False
        x = base_model(x)
        x = Conv2D(32, (3, 3), activation='relu', kernel_initializer='he_normal', padding='same')(x)
        x = MaxPooling2D((2, 2))(x)

        x = Flatten()(x)
        x = Dense(128, activation='relu', kernel_initializer='he_uniform')(x)

        self.outputs = Dense(1, activation='sigmoid')(x)

        self.model = Model(inputs=self.inputs, outputs=self.outputs, name=self.name)

        self.model.compile(optimizer=optimizer, loss=loss, metrics=metrics, loss_weights=loss_weights)

        if show_summary:
            self.model.summary()

```

```
model = SuperCatDogClassifierV2(img_height, img_width, img_channels, classes, training=True, gpu_initialized=True)
```

```
model.compile(optimizer=Adam(lr = 1e-4), loss='binary_crossentropy', metrics=['accuracy'],
show_summary=True)

# current time
current_time = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")

# create logging
log_dir = os.path.join(project_path, f'logs\{model.name}\{current_time}')

# create all callbacks
callbacks = [
    EarlyStopping(patience=10, monitor='val_loss'),
    TensorBoard(log_dir=log_dir, profile_batch=0)
]
```

```

WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.iter
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.iter
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_1
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_1
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_2
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.beta_2
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.decay
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.decay
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.learning
_rate
WARNING:tensorflow:Unresolved object in checkpoint: (root).optimizer.learning
_rate
WARNING:tensorflow:A checkpoint was restored (e.g. tf.train.Checkpoint.restore
or tf.keras.Model.load_weights) but not all checkpointed values were used.
See above for specific issues. Use expect_partial() on the load status objec
t, e.g. tf.train.Checkpoint.restore(...).expect_partial(), to silence these w
arnings, or use assert_consumed() to make the check explicit. See https://ww
w.tensorflow.org/guide/checkpoint#loading_mechanics for details.
WARNING:tensorflow:A checkpoint was restored (e.g. tf.train.Checkpoint.restore
or tf.keras.Model.load_weights) but not all checkpointed values were used.
See above for specific issues. Use expect_partial() on the load status objec
t, e.g. tf.train.Checkpoint.restore(...).expect_partial(), to silence these w
arnings, or use assert_consumed() to make the check explicit. See https://ww
w.tensorflow.org/guide/checkpoint#loading_mechanics for details.
Model: 'SuperCatDogClassifierV2'

```

Layer (type)	Output Shape	Param #
input_3 (InputLayer)	[(None, 224, 224, 3)]	0
rescaling (Rescaling)	(None, 224, 224, 3)	0
mobilenetv2_1.00_224 (Functi	(None, 7, 7, 1280)	2257984
conv2d_6 (Conv2D)	(None, 7, 7, 32)	368672
max_pooling2d_6 (MaxPooling2	(None, 3, 3, 32)	0
flatten_2 (Flatten)	(None, 288)	0
dense_4 (Dense)	(None, 128)	36992
dense_5 (Dense)	(None, 1)	129
Total params: 2,663,777		
Trainable params: 405,793		
Non-trainable params: 2,257,984		

```
# fit the model using the training data
```

```
results2 = model.fit(
    training=dataset.train_ds,
    callbacks=callbacks,
    epochs=epochs,
    validation=dataset.val_ds,
    validation_steps=dataset.val_size,
    steps_per_epoch=dataset.train_size)
```

```
# save the weights of the model
```

```
weights_path = os.path.join(project_path, f'models\{model.name}_trained_model_weights')
model.save_weights(weights_path)
```

Epoch 1/100

218/218 [=====] - 15s 67ms/step - loss: 0.0654 - accuracy: 0.9757 - val\_loss: 0.0314 - val\_accuracy: 0.9881

Epoch 2/100

218/218 [=====] - 15s 71ms/step - loss: 0.0150 - accuracy: 0.9956 - val\_loss: 0.0307 - val\_accuracy: 0.9896

Epoch 3/100

218/218 [=====] - 16s 75ms/step - loss: 0.0058 - accuracy: 0.9990 - val\_loss: 0.0297 - val\_accuracy: 0.9894

Epoch 4/100

218/218 [=====] - 17s 76ms/step - loss: 0.0021 - accuracy: 0.9998 - val\_loss: 0.0302 - val\_accuracy: 0.9896

Epoch 5/100

218/218 [=====] - 16s 74ms/step - loss: 9.6561e-04 - accuracy: 1.0000 - val\_loss: 0.0319 - val\_accuracy: 0.9907

Epoch 6/100

218/218 [=====] - 14s 65ms/step - loss: 3.8996e-04 - accuracy: 1.0000 - val\_loss: 0.0330 - val\_accuracy: 0.9907

Epoch 7/100

218/218 [=====] - 15s 68ms/step - loss: 2.3281e-04 - accuracy: 1.0000 - val\_loss: 0.0336 - val\_accuracy: 0.9907

Epoch 8/100

218/218 [=====] - 16s 75ms/step - loss: 1.7008e-04 - accuracy: 1.0000 - val\_loss: 0.0343 - val\_accuracy: 0.9907

Epoch 9/100

218/218 [=====] - 16s 73ms/step - loss: 1.3201e-04 - accuracy: 1.0000 - val\_loss: 0.0348 - val\_accuracy: 0.9905

Epoch 10/100

218/218 [=====] - 16s 73ms/step - loss: 1.0583e-04 - accuracy: 1.0000 - val\_loss: 0.0353 - val\_accuracy: 0.9905

Epoch 11/100

218/218 [=====] - 16s 73ms/step - loss: 8.6372e-05 - accuracy: 1.0000 - val\_loss: 0.0359 - val\_accuracy: 0.9902

Epoch 12/100

218/218 [=====] - 16s 73ms/step - loss: 7.1457e-05 - accuracy: 1.0000 - val\_loss: 0.0364 - val\_accuracy: 0.9902

Epoch 13/100

218/218 [=====] - 16s 73ms/step - loss: 5.9817e-05 - accuracy: 1.0000 - val\_loss: 0.0370 - val\_accuracy: 0.9900

```
# re initialize the model.
model.training = False
model.compile(optimizer=Adam(lr = 1e-4), loss='binary_crossentropy', metrics=['accuracy'],
show_summary=False)
model.load_weights(weights_path)

print('\n# Evaluate on test data')
result = model.evaluate(dataset.actual_test_ds)
print('test loss, test acc:', result)
res = dict(zip(model.get_metric_names(), result))
print(res)
```

```
# Evaluate on test data
```

```
62/62 [=====] - 6s 96ms/step - loss: 0.0452 - accuracy: 0.9892
test loss, test acc: [0.04516126587986946, 0.9891632795333862]
{&#39;loss&#39;: 0.04516126587986946, &#39;accuracy&#39;: 0.9891632795333862}
```

```
y_pred = model.predict((image_batch, label_batch))
```

```
1/2 [=====&gt;.....] - ETA: 0sWARNING:tensorflow:Callbacks
method `on_predict_batch_end` is slow compared to the batch time (batch time:
0.0070s vs `on_predict_batch_end` time: 0.0210s). Check your callbacks.
WARNING:tensorflow:Callbacks method `on_predict_batch_end` is slow compared t
o the batch time (batch time: 0.0070s vs `on_predict_batch_end` time: 0.0210
s). Check your callbacks.
2/2 [=====] - 0s 15ms/step
```

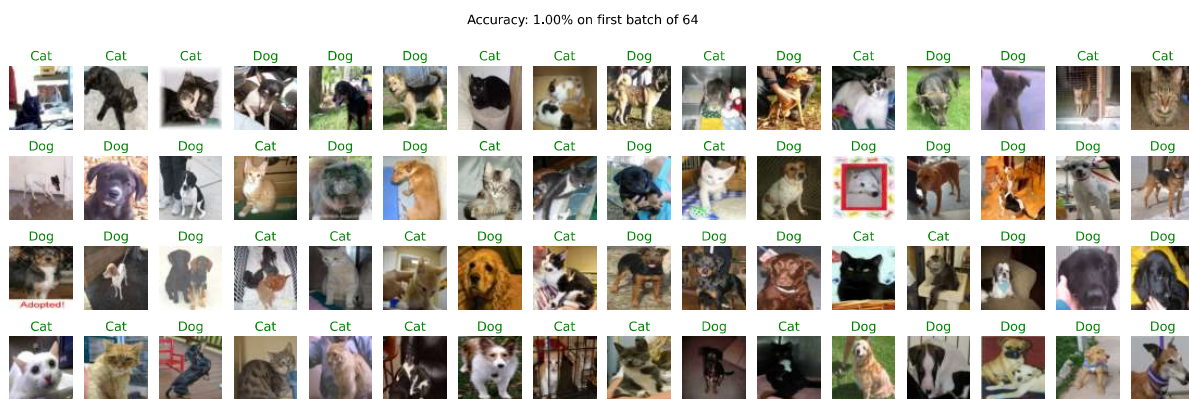


```

predictions = [round(yhat[0]) for yhat in y_pred]
test_accuracy = sum(predictions == label_batch.numpy()) / len(predictions)
fig = plt.figure(figsize=(20, 6))
for i in range(batch_size):
    ax = plt.subplot(4, 16, i + 1)
    plt.imshow(image_batch[i])
    label = label_batch[i]
    pred = predictions[i]
    color = 'r'
    if label == pred:
        color = 'g'
    fontdict = { 'color': color }
    plt.title(class_names[label], fontdict = fontdict)
    plt.axis("off")
fig.suptitle(f'Accuracy: {test_accuracy:.2f}% on first batch of {batch_size}')

```

Text(0.5, 0.98, &#39;Accuracy: 1.00% on first batch of 64&#39;)



## Results

The results of the model using a MobileNetV2 as base model are impressive. It achieved 98% accuracy on the test set. I do think that it is able to achieve this accuracy because of the weights that are used. ImageNet contains animal classes, so the base model probably has been trained on cat and dogs as well.

```
acc1 = results1.history['accuracy']
val_acc1 = results1.history['val_accuracy']

loss1 = results1.history['loss']
val_loss1 = results1.history['val_loss']

acc2 = results2.history['accuracy']
val_acc2 = results2.history['val_accuracy']

loss2 = results2.history['loss']
val_loss2 = results2.history['val_loss']

plt.figure(figsize=(8, 8))
plt.subplot(2, 1, 1)
plt.plot(acc1, label='V1 Training Accuracy')
plt.plot(val_acc1, label='V1 Validation Accuracy')
plt.plot(acc2, label='V2 Training Accuracy')
plt.plot(val_acc2, label='V2 Validation Accuracy')
plt.legend(loc='lower right')
plt.ylabel('Accuracy')
plt.ylim([min(plt.ylim()),1])
plt.title('Training and Validation Accuracy')

plt.subplot(2, 1, 2)
plt.plot(loss1, label='V1 Training Loss')
plt.plot(val_loss1, label='V1 Validation Loss')
plt.plot(loss2, label='V2 Training Loss')
plt.plot(val_loss2, label='V2 Validation Loss')
plt.legend(loc='upper right')
plt.ylabel('Cross Entropy')
plt.ylim([0,1.0])
plt.title('Training and Validation Loss')
plt.xlabel('epoch')
plt.show()
```

